

Speech Recognition: Key Word Spotting through Image Recognition



Alex Spaeth, David Harrison, Peter Braun, Salil Kanetkar, Sanjay
Krishna Gouda

Objective


- Keyword spotting (or more simply, word spotting) is a problem that was historically first defined in the context of speech processing .
- In speech processing, keyword spotting deals with the identification of pre defined keywords in utterances.
- In this project, we will work on the task of voice command recognition, following the TensorFlow Speech Recognition Challenge that is being carried out on Kaggle.

Dataset

- Speech Commands Dataset
- 65,000 labeled audio examples
 - 1 second each
- Contains 41 different words
- 12 target labels
 - E.g. “up”, “stop”, “down”...
 - Also “<silence>” and “<unknown>”


Speech to Image

- Tried 1-D CNN on raw audio with bad results, so tried to apply 2-D CNNs.
- Convert audio into meaningful images using spectrograms (a SciPy function).
- Plots of frequency content vs. time.
- Each column is the Fourier transform of the input in a window of time.
- Promising because different words actually look different! →



Speech to Image

- Use spectrograms because a plot of amplitude vs. time is uninformative →
- Can only recognize volume envelope.
- Don't expect good performance, but we did try this type of image in one model for comparison.
- Spectrograms structure the information better, allowing analysis in terms of pitch and time rather than absolute sound pressure.



Models - Initial


a CNN with two pooling convolutional layers and one FC layer.

- the model suggested by the authors of the challenge
- a large model: about 800 million FLOPs per inference, and 940,000 weights
- good performance (89% validation), but convergence takes > 8h
- compared some more computationally tractable models

Models

Low latency convolutional Model:

- Convolution Layer:
 - Filter: 7x7
 - Stride: 3
- Dnn1: 100 units (no nonlinearity)
- dnn2 : 100 units (relu)
- 63,800 Parameters
- Estimated: 11 million FLOPs (7 times less)



Results

Low Latency Convolutional Model

Xavier Initialization:


- Matrix \mathbf{M} (n by m)
- Assign values uniformly sampled from $[-\epsilon, \epsilon]$
- $\epsilon = \sqrt{6} / \sqrt{(m + n)}$

Best Accuracies:

- Xavier: 87.6%, epoch 196
- Truncated Normal: 78.7%, epoch 200

We also tried constant initialization methods

- all weights = 0.5
- did not converge




Results

Low Latency Convolutional Model

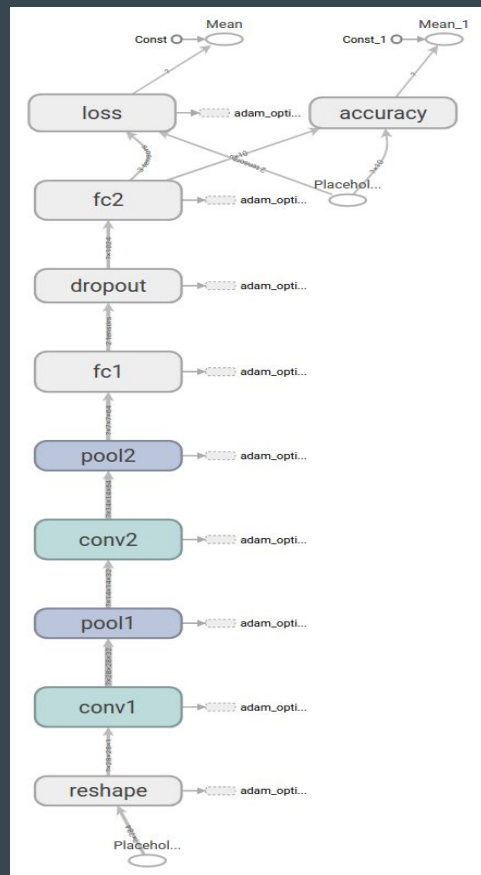
Optimization Algorithms:

- Adam: adaptive learning rate
 - Best accuracy, epoch: 87.1%, 52
- SGD
 - Best accuracy, epoch: 87.6%, 196



Tensorflow MNIST expert

- We also tested the deep CNN available on the Tensorflow website
- It has two convolutional layers, 2 pooling layers and a fully connected layer
- It also employs dropout to combat overfitting and uses ReLu as the activation function.



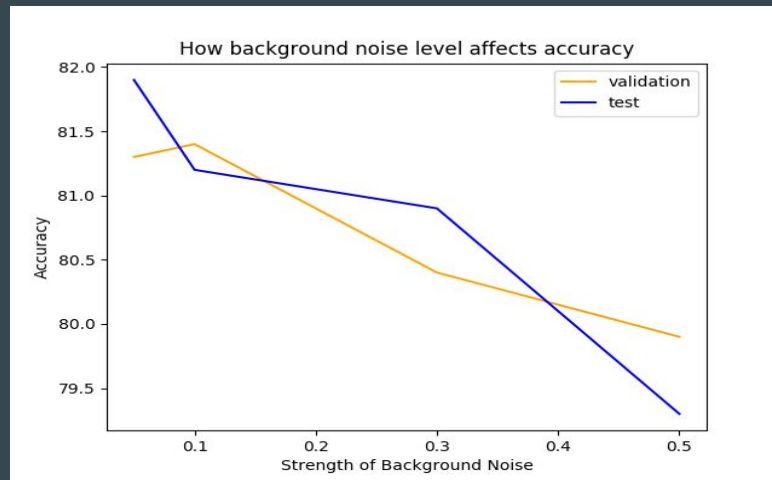
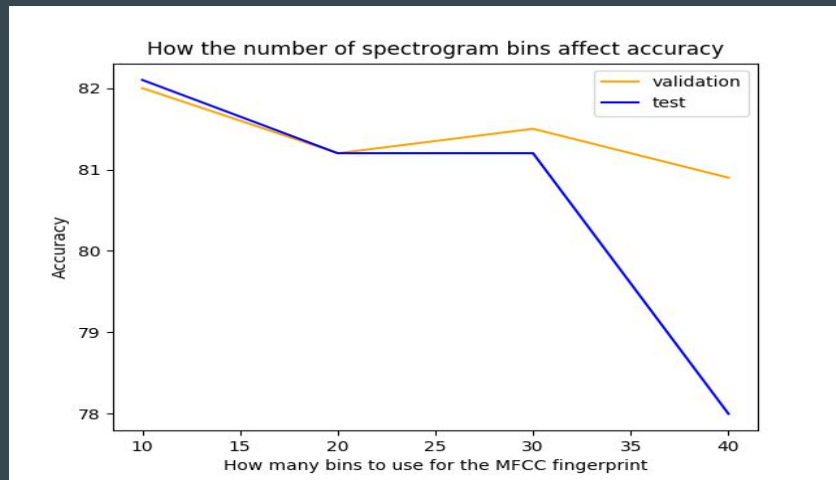
Models

MNIST CNN

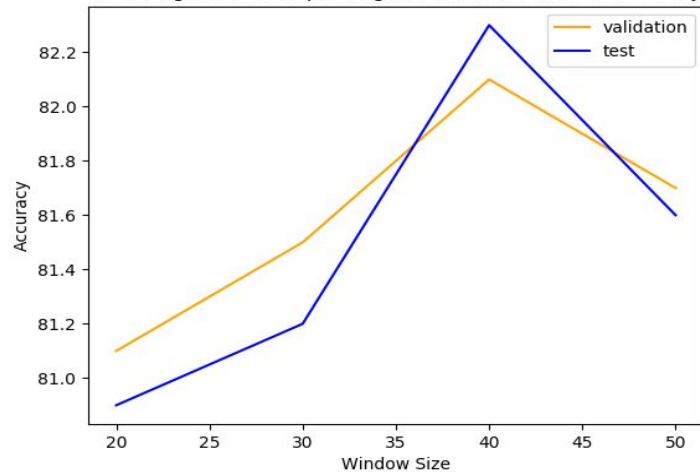
- We used the MNIST sample CNN architecture for this process.
- It is a simple, neural network with two convolutional layers and two maxpool layers. The architecture is similar to the default model, but much smaller.

Results

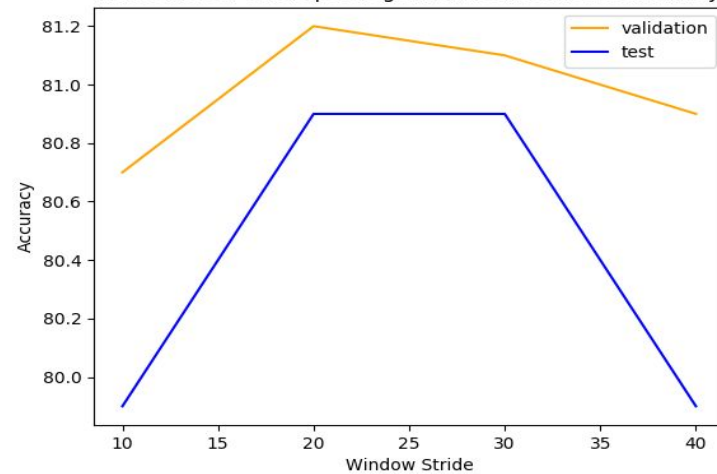
- Varying hyperparameters on Low Latency Convolutional Model



How length of each spectrogram timeslice affects accuracy



How stride of each spectrogram timeslice affects accuracy



also tried a CNN intended for MNIST handwriting classification

- on log spectrograms of the audio files: validation accuracy 52.7%
- on time vs. amplitude images: validation accuracy 50.3%

Model - CNN

A 5-layer CNN and a 7-layer CNN.

Two variations of the bigger architecture. Different number of filters.

Two training variants - vanilla CNN and CNN with adversarial training.

Architecture:

- Group of convolutions using Convolutional layer followed by max pooling layer.
- Flatten.
- Fully Connected.
- Dropout? Further.

CNN - Architecture

Idea from MCDNN.

(#2 for MNIST)

Smaller CNN :

Input \rightarrow (convolutional layer \rightarrow relu+dropout \rightarrow max pooling) $\times 3 \rightarrow$ flatten \rightarrow fully connected \rightarrow fully connected + dropout.

Larger CNN :

Input \rightarrow (convolutional layer \rightarrow elu \rightarrow max pooling) $\times 5 \rightarrow$ fully connected \rightarrow fully connected \rightarrow optional dropout.

Also differ in number of filters used. Details in report (too many to include in here).

Adversarial Training

Followed Szegedy et al [2012] and
Mikolov (VAT) to add adversarial training.

Method:


Augment more data (images) that are
“visually similar” to existing ones.

Why:

Adds regularization.


```
plt.imshow(train_x[12])
```

```
Out[4]: <matplotlib.image.AxesImage at 0x1fe52919160>
```



```
In [8]: train_x2 = np.add(train_x*0.01*np.std(train_x),train_x)  
plt.imshow(train_x2[12])
```

```
Out[8]: <matplotlib.image.AxesImage at 0x1fe2d80fd30>
```



```
In [11]: np.array_equal(train_x[12],train_x2[12])
```

```
Out[11]: False
```

Architecture:

7 layers

input -> (conv_2d -> elu -> max_pool ->) x 5 --> flatten --> fully connected --> fully connected

Details:

Conv1 : 7x7 filters x8 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 8, strides: 8,8

Conv2 : 5x5 filetrs x16 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 4, strides: 2,2

Conv3 : 3x3 filters x8 (1,1) strides followed by elu transformation and max pooling with n_h = n_w = 2, strides: 2,2

Conv3 : 2x2 filters x4 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 4, strides: 2,2

Conv3 : 2x2 filters x4 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 2, strides: 2,2

Note: For this setting, added dropout regularization at all the elu transformation layers and at the fully connected layers but it seemed like the models were underfitting (over regularization) so not including those results.

Larger Architecutre.

Architecture:

7 layers

input -> (conv_2d -> elu -> max_pool ->) x 5 --> flatten --> fully connected --> fully connected

Details:

Conv1 : 7x7 filters x64 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 8, strides: 8,8

Conv2 : 5x5 filetrs x32 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 4, strides: 2,2

Conv3 : 3x3 filters x16 (1,1) strides followed by elu transformation and max pooling with n_h = n_w = 2, strides: 2,2

Conv3 : 2x2 filters x8 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 4, strides: 2,2

Conv3 : 2x2 filters x4 (2,2) strides followed by elu transformation and max pooling with n_h = n_w = 2, strides: 2,2

img res	batch_size	n_epochs	l_rate	tr_acc	te_acc	activation	loss
28_28	1024	5000	0.001	0.877098	0.86575	elu	0.38977
288_288	1024	5000	0.001	0.878353	0.867295	elu	0.39376
28_28 + VAT	1024	5000	0.001	0.719849	0.868531	elu	0.37929
288_288 + VAT	1024	5000	0.001	0.718285	0.860188	elu	0.38031


Run 1 thres <= 0.1

img res	batch_size	n_epochs	l_rate	tr_acc	te_acc	activation	loss
28_28	1024	168(e)	0.009	0.973021	0.926078	elu	0.0906
288_288	1024	171(e)	0.009	0.965238	0.918894	elu	0.0942
28_28 + VAT	1024	105(e)	0.009	0.969892	0.921984	elu	0.0954
288_288 + VAT	1024	97 (e)	0.009	0.960991	0.915109	elu	0.0978


Run 2 thres <= 0.05

img res	batch_size	n_epochs	l_rate	tr_acc	te_acc	activation	loss
28_28	1024	411(e)	0.009	0.986115	0.921675	elu	0.0539
288_288	1024	440(e)	0.009	0.983778	0.922756	elu	0.0535
28_28 + VAT	1024	243(e)	0.009	0.981424	0.918198	elu	0.0549
288_288 + VAT	1024	236(e)	0.009	0.977743	0.915572	elu	0.0551


training vs validation



costs



training vs validation



costs

